



Mirror's Edge: Policy Gradient

AI Parkour Agent

Victor Tenneroni | vt2435@columbia.edu

Problem Statement

Problem: Navigate randomized 3D platforms with resource constraints

- 20 platforms, gaps 2.5-4.5 units, widths 20-84 units
- Stamina system (100 max): Sprint costs 20/s, Roll costs 60, Jump costs 20

Why RL?

- High-dimensional state (14 obs) + complex action space (5 discrete)
- Infinite environment variations → must generalize, not memorize
- Strategic tradeoffs: speed vs stamina conservation
- No closed-form solution for stamina dynamics + randomization


Traditional Approaches Fail

- Rule-based: Cannot handle randomization
- PID control: No strategic resource management
- Fixed environments: Agent memorizes sequences



<https://youtu.be/2PWbs4OHtTM>

MDP



Algo: PPO with high exploration

Why PPO?

- vs DQN: Better for continuous observations (14-dim)
- vs A2C: Clipped objective prevents destructive updates
- vs TRPO: Simpler, computationally efficient

Key Hyperparameters:		
Parameter	Value	Rationale
Learning rate	3e-4 (decay)	Standard PPO
Batch size	1024	Balance efficiency/noise
β (entropy)	0.1 → 0.00074	High exploration (6.7x test_v1)
ϵ (clipping)	0.2	Prevent large updates
λ (GAE)	0.95	Bias-variance tradeoff
Epochs	5	Multiple passes

Network Architecture:

- Actor-critic
- Policy: 2×256 hidden → 5 action logits
- Value: 2×128 hidden (separate network)

Tweaks:

1. **Episodic style bonus** (40% episodes): +1.5 roll reward
2. **Dual roll rewards**: +0.5 base (always) + 1.5 style (conditional)
3. **Low stamina penalty**: -0.002/step when <20%
4. **Action constraints**: Environment blocks invalid actions

Iterative Reward Design

Goal Alignment: Forward progress (primary) + time efficiency + stamina management + strategic risk-taking

Iterative Reward Design

- test_v11 (baseline): +0.1 progress, +10 target → +11.15 reward, no roll usage
- =run28 (sprint added): 12 u/s, 33.33/s cost → +78.32 reward (+603%), but sprint bashing (38% usage, stamina=0)
- training_171550 (roll v1): Cost 150, bonus +0.1 (15% eps) → +67.90 reward, rolls rarely used (7.5s regen too slow)
- training_210205 (final): Cost 60, base +0.5 (always), style +1.5 (40% eps) → +89.18 reward (+31%), rolls significantly increased

Key Improvements

1. Dual Roll Rewards: Base (+0.5 always) prevents learned aversion, style (+1.5 episodic) creates exploration
2. Multi-timescale Structure: Dense (progress, time, stamina) guide learning, sparse (target, fall) define success
3. Balanced Penalties: Time -0.001 (efficiency), low stamina -0.002 (conservation), fall -1.0 (moderate, not catastrophic)
4. Stamina Rebalancing: Sprint 20/s (reduced), regen 30/s (increased) → allows buildup for rolls/jumps

State Space Design

Critical: Platform raycasts [2f, 4f, 6f, 8f, 10f] - test_v9 (no raycasts) +3.43 vs test_v10 (raycasts) +9.85 (187% improvement)

Encoding: Target position (raw \mathbb{R}^3), velocity (raw \mathbb{R}^3), stamina [0,1], raycasts [0,1], grounded {0,1}

Why raycasts essential: Randomized environments prevent memorization → agent must perceive terrain

Reward structure drives strategy: Per-unit-time rewards → sprint bashing.

Solution: Cost-benefit balance via stamina penalties + achievable action costs + dual reward structure for underutilized actions.

Results - 30th training loop

Final Reward: +89.18 (at 2M steps)

Previous Best: +78.32 (run28)

Improvement: +14% over previous best

Training Progression:

- 500k steps: +26.67
- 1.0M steps: +45.25
- 1.5M steps: +81.60
- 2.0M steps: +89.18

Policy Loss: 0.0233 (mean, range 0.0175-0.0312)

Value Loss: 0.985 (mean, range 0.400-1.808)

Policy Entropy: 0.657 (mean, range 0.657-1.605)

Learning Rate: 8.36e-07 (final, decayed from 3.0e-4)

Epsilon: 0.100 (final, decayed from 0.2)

Beta: 0.000289 (final, decayed from 0.1)

Episode Statistics

- Mean Episode Reward: 80.06 (range 3.05-88.82)
- Mean Episode Length: 61.07 steps (range 4.90-68.50)
- Mean Max Distance: 555.91 units (range 29.89-603.56)
- Mean Episode Duration: 609.64 environment steps

Action Distribution

- Jog: 67.61% (primary movement)
- Sprint: 14.00%
- Roll: 7.81% (increased from 0.69% in previous run)
- Jump: 3.53%
- Idle: 7.04%

Action Counts (per episode, mean):

- Jog: 2072 actions
- Sprint: 424 actions
- Roll: 239 actions
- Jump: 102 actions
- Idle: 216 actions

Agent Behavior

- Roll Usage: 7.81% of actions (vs 0.69% previous run)
- Roll Count: 239 rolls per episode (mean)
- Roll Improvement: 11.3x increase over previous run

Critical Experiments & Lessons over 30 training loops

Experiment 1: Axis Bug (test_v4)

- **Setup:** Progress reward tracking Z-axis instead of X-axis
- **Result:** -2.497 reward, agent learned to idle (safest strategy)
- **Lesson:** Always verify reward signals track correct objective

Experiment 2: Training Duration (test_v5 vs test_v6)

- **Setup:** 500k steps vs 2M steps, same config
- **Result:** +5.976 vs +8.478 reward (+42% improvement)
- **Lesson:** 500k insufficient for convergence in this domain

Experiment 3: Randomization Requires Perception (test_v9 vs test_v10)

- **Setup:** test_v9 (no raycasts, randomized platforms) vs test_v10 (5 raycasts, randomized)
- **Result:** +3.43 vs +9.85 reward (187% improvement, 60% drop from fixed environment)
- **Lesson:** Fixed environments allow memorization, random environments require sensors
- **Key insight:** Agent without raycasts learned "run 50 steps, jump, repeat" in fixed layout, completely failed in randomized layout

Experiment 4: Beta Coefficient Impact (test_v11 vs =run28)

- **Setup:** $\beta=0.015$ (low entropy) vs $\beta=0.1$ (high entropy)
- **Result:** +11.15 vs +78.32 reward (603% improvement)
- **Lesson:** Complex action spaces need high exploration to discover optimal strategies
- **Evidence:** test_v11 entropy=0.035 (converged to local optimum), =run28 entropy=0.597 (diverse strategies)

Experiment 5: Roll Cost Tuning (training_20251207_171550 vs training_20251207_210205)


- **Setup:** Roll cost 150 ($1.5 \times$ max) vs 60 ($0.6 \times$ max)
- **Result:** +67.90 vs +89.18 reward (+31% improvement)
- **Lesson:** Action costs must be achievable mid-episode, not just theoretically available
- **Evidence:** Roll cost 150 requires 7.5s regeneration from 0 stamina (unrealistic during gameplay)

Demo (inference)



<https://youtu.be/geAkthimal0>

Architecture & Stack



Future Work

SHORT-TERM IMPROVEMENTS

1. Hyperparameter Optimization

- Current: Beta 0.1 (linear), entropy 0.635 (over-exploring), LR 3e-4
- Proposed: Beta 0.05→0.001, LR 5e-4→1e-5, Epsilon 0.15→0.05 (all exponential), Lambda 0.98, Epochs 3
- Rationale: Exponential decay matches learning curves; current linear too slow
- Expected: +95-100 reward (vs 89.18), faster convergence, lower entropy ~0.2-0.3

2. Movement Smoothing

- Issue: Sprint stuttering, no flow penalty
- Solution: Momentum system - continuous sprint 12→14 u/s, interruption penalty -0.005

LONG-TERM DIRECTIONS

1. Full 3D Movement - Multi-axis platforms, turning

2. Expanded Action Space - Add slide, wall jump, vault (5→9+ actions or continuous control)

3. (Dynamic) Obstacles - Moving platforms, time-dependent physics, partial observability



4. LLM-Assisted Tuning - Automated hyperparameter optimization on environment changes

5. Q-Learning, DQN Benchmark

Thank you!

Appendix

Scaling Training in Unity



Sprint Bashing Demo



Prior Open Questions & Initial Development Path

Core Problems:

- Continuous → discrete state space conversion
- 2D → 3D navigation scaling
- Vision-based vs positional observations
- Fast-forward training in Unity

Two Development Tracks:

- Vertical: Add actions (slide, salto, wall-jump)
- Horizontal: Compare RL methods (PPO → Q-Learning → DQN)

Success = Agent navigates A→B with visible human influence on behavior

MDP

State Space $S \subseteq \mathbb{R}^{14}$:

- Target relative position (3): (Δx , Δy , Δz)
- Velocity (3): (v_x , v_y , v_z) units/sec
- **Platform raycasts (5)**: Downward at [2f, 4f, 6f, 8f, 10f], normalized [0,1]
 - **Critical**: 60% performance drop without raycasts
- Obstacle distance (1): Forward raycast [0,1]
- Grounded (1): Binary {0,1}
- Stamina (1): Normalized [0,1]

Reward Function:

- Dense: $+0.1 \times \text{progress}$, $+0.001 \times \text{grounded}$, $-0.001 \times \text{time}$, $-0.002 \times \text{low stamina}$
- Sparse: $+10.0 \times \text{target}$, $-1.0 \times \text{fall}$
- Roll: $+0.5 \times \text{base (always)}$ + $1.5 \times \text{style (40\% episodes)}$


Action Space $A = \{0,1,2,3,4\}$:

Action	Speed	Stamina Cost	Constraints
Idle	0	0	Always
Jump	Impulse	20	Ground + stamina ≥ 20
Jog	6	0	Always
Sprint	12	20/s	Stamina > 0 + 0.5s cooldown
Roll	18	60	Stamina ≥ 60 + not rolling

Past Experience making Games with AI



Raycasting engine in C with
Stable Diffusion Textures



LLM RPG in BabylonJS

<https://drive.google.com/file/d/110FcmJgdU1eOwg5Pv-QXF9nvBX9XceBS/view?resourcekey>

Policy Gradient Learning System

Why PPO:

- Trial-and-error learning (no expert demos)
- Handles delayed consequences
- Proven for continuous control in Unity ML-Agents

Stack: Unity 2022.3 + ML-Agents + Python + Cloud GPU

